EMPOWERING MRI-BASED BRAIN TUMOR DIAGNOSIS WITH DEEP CNN'S FOR ACCURATE DETECTION, INTELLIGENT STAGING, AND AUTOMATED CLINICAL REPORTING

Submitted in partial fulfillment of the requirements for the award of Bachelor of Engineering degree in Computer Science and Engineering

By

YOGARAJAN N (Reg. No- 42111505)

YOKESHMADAN M (Reg. No- 42111510)

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

SCHOOL OF COMPUTING

SATHYABAMA

INSTITUTE OF SCIENCE AND TECHNOLOGY
(DEEMED TO BE UNIVERSITY)
CATEGORY - 1 UNIVERSITY BY UGC
Accredited with Grade "A++" by NAAC I Approved by AICTE
JEPPIAAR NAGAR, RAJIV GANDHI SALAI, CHENNAI – 600119

APRIL - 2026

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

BONAFIDE CERTIFICATE

This is to certify that this Project Report is the bonafide work of **YOKESHMADAN M(42111510)**, **YOGARAJAN N (42111505),** who carried out the Project entitled "**EMPOWERING MRI-BASED BRAIN TUMOR DIAGNOSIS WITH DEEP CNN'S FOR ACCURATE DETECTION, INTELLIGENT STAGING, AND AUTOMATED CLINICAL REPORTING**" under my supervision from November 2025 to April 2026.

Internal Guide

Ms. P. KRISHNAVENI, M-Tech., (Ph.D).,

Head of the Department

Dr. L. LAKSHMANAN, M.E., Ph.D.,

Submitted for Viva voce Examination held on_____

Internal Examiner **External Examiner**

DECLARATION

I, **YOKESHMADAN M(Reg. No- 42111510),** hereby declare that the Project Report entitled **"EMPOWERING MRI-BASED BRAIN TUMOR DIAGNOSIS WITH DEEP CNN'S FOR ACCURATE DETECTION, INTELLIGENT STAGING, AND AUTOMATED CLINICAL REPORTING"** done by me under the guidance of **Ms. P. KRISHNAVENI, M-Tech., (Ph.D).,**is submitted in partial fulfillment of the requirements for the award of Bachelor of Engineering degree in **Computer Science and Engineering**.

DATE:

PLACE: Chennai **SIGNATURE OF THE CANDIDATE**

ACKNOWLEDGEMENT

I am pleased to acknowledge my sincere thanks to **Board of Management of SATHYABAMA INSTITUTE OF SCIENCE AND TECHNOLOGY** for their kind encouragement in doing this project and for completing it successfully. I am grateful to them.

I convey my thanks to **Dr. L. LAKSHMANAN, M.E., Ph. D., DEAN AND HEAD**, School of Computing, Department of Computer Science and Engineering for providing me necessary support and details at the right time during the progressive reviews.

I would like to express my sincere and deep sense of gratitude to my Project Guide **Ms. P. KRISHNAVENI, M-Tech., (Ph.D).,** for her valuable guidance, suggestions and constant encouragement paved way for the successful completion of my project work.

I wish to express my thanks to all Teaching and Non-teaching staff members of the **DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING** who were helpful in many ways for the completion of the project.

ABSTRACT

Brain tumor detection and classification present critical challenges in medical imaging due to complex tumor morphology, size variability, and the high stakes of accurate diagnosis. Conventional approaches rely on manual MRI analysis, which is time-consuming and prone to human error. This research proposes an automated deep learning system using Python, TensorFlow, Keras, and OpenCV, following a multi-stage pipeline that includes MRI preprocessing (normalization, noise reduction, skull stripping, and data augmentation), U-Net-based tumor segmentation, and CNN classification of tumor types — glioma, meningioma, and pituitary tumors.

The system automates tumor detection from raw MRI scans, segments affected regions accurately, and performs stage-wise classification based on tumor characteristics. Data augmentation improves model generalization, while segmentation visualizations enhance interpretability for medical professionals. Performance is evaluated using accuracy, sensitivity, specificity, and F1-score. By minimizing diagnostic delays and reducing dependency on manual interpretation, the system assists radiologists in clinical decision-making, supports personalized treatment planning, and offers a scalable solution for healthcare facilities — contributing to improved patient outcomes through timely and accurate diagnosis.

TABLE OF CONTENTS

CHAPTER NO	TITLE	PAGE NO
	ABSTRACT	**V**
	CONTENTS	**VI**
	LIST OF ABBREVIATION	**VIII**
	LIST OF FIGURE	**IX**
1	**INTRODUCTION**	**1**
	1.1 Background	1
	1.2 Problem Statement	3
	1.3 Objectives	4
	1.4 Scope of the Project	5
	1.5 Significance of the Study	6
2	**LITERATURE SURVEY**	**8**
	2.1 Literature Survey	8
	2.2 Inferences and Challenges in Existing Systems	14
3	**AIM AND SCOPE OF THE PRESENT INVESTIGATION**	**18**
	3.1 Aim of the Study	18
	3.2 Scope of the Study	22
4	**EXPERIMENTAL MATERIALS, METHODS, AND ALGORITHMS USED**	**28**
	4.1 Dataset Description, and Data Engineering	28

	4.2 Advanced Image Preprocessing and Fuzzy-Logic Segmentation	28
	4.3 System Architecture and Deep CNN Implementation	29
	4.4 Mathematical Optimization and Training Hyperparameters	32
	4.5 System Integration, Gradio Deployment, and Clinical Reporting	32
5	**RESULTS AND DISCUSSION, PERFORMANCE ANALYSIS**	**34**
	5.1 Classification Performance Evaluation	34
	5.2 Comparative Analysis with Existing Models	35
	5.3 Real-Time Performance and Inference Analysis	37
	5.4 Discussion and Critical Analysis	38
6	**SUMMARY AND CONCLUSIONS**	**40**
	6.1 Summary of the Study	40
	6.2 Conclusions and Future Work	43
	REFERENCES	**45**
	ANNEXURES	**46**

LIST OF ABBREVIATION

ABBREVIATION	FULL FORM
Adam	Adaptive Moment Estimation
AI	Artificial Intelligence
ANN	Artificial Neural Networks
CAD	Computer-Aided Diagnosis
CapsNet	Capsule Networks
CNN	Convolutional Neural Networks
CT	Computed Tomography
DICIM	Digital Imaging and Communications in Medicine
DNN	Deep Neural Network
DWT	Discrete Wavelet Transform
FCM	Fuzzy C-Means
GLCM	Gray Level Co-occurrence Matrix
GPU	Graphics Processing Unit
KELM	Kernel Extreme Learning Machines
KNN	K-Nearest Neighbors
LBP	Local Binary Pattern
MRI	Magnetic Resonance Imaging
PACS	Picture Archiving and Communication Systems
PCA	Principal Component Analysis
PET	Positron Emission Tomography
SVM	Support Vector Machines
VGG16	Visual Geometry Group 16

LIST OF FIGURES

FIGURE NO	TITLE	PAGE NO
4.2.1	System Architecture	30
4.5.1	Website Interface	33
5.1.1	Training and Validation Accuracy and Loss Curves	35

CHAPTER 1
INTRODUCTION

1.1 Background

When it comes to brain tumors, time is everything. These growths — whether benign or malignant — build up pressure inside the skull that can permanently damage a person's ability to think, move, or even survive if they aren't caught early enough. The brain is an unforgiving environment in this sense; unlike many other organs, it has virtually no room to accommodate an expanding mass, and the consequences of delayed detection can escalate from manageable to catastrophic within a matter of weeks. MRI scans remain our best non-invasive window into the brain, offering a level of soft tissue detail that no other widely available imaging modality can match — but the sheer volume of scans produced every day in hospitals far exceeds what radiologists can realistically review with the care and attention each one deserves. Reading an MRI is also deeply subjective — tumor edges can be blurry, tissue differences are subtle, and even the most experienced specialist can miss something when they're tired or working under the kind of relentless pressure that modern radiology departments routinely face. The variability in how different radiologists interpret the same scan is well-documented in the medical literature, and while experienced clinicians generally perform remarkably well, the margin for error in a domain this consequential is uncomfortably wide. Compounding this is the growing global shortage of specialist radiologists, particularly in lower-income countries and rural regions, where patients may wait days or weeks for a report that a well-resourced urban hospital might produce in hours. The result is a system where delays between scanning and diagnosis can quite literally cost lives, and where the gap between the diagnostic support that patients need and what the healthcare system can reliably deliver continues to widen.

This is precisely why there is a growing push for smarter, more automated diagnostic tools — not to replace the radiologist, but to give them better support and to extend reliable diagnostic capability into settings where specialist access is limited. Most of today's computer-aided diagnosis systems fall short of what is genuinely needed; they struggle with how wildly different tumors can look from one patient to the next, and they often stop at detection without helping doctors understand what comes

next or how urgently they need to act. This project takes a meaningfully different approach by building a system that treats detection as the beginning of the diagnostic process rather than the end of it. By using fuzzy thresholding for segmentation, the system handles the natural blurriness at tumor boundaries far more effectively than traditional binary methods, allowing for a much more accurate measurement of how large a tumor actually is relative to the surrounding brain tissue — something that matters enormously when deciding how serious a case is and how aggressively it needs to be treated. The difference between a tumor that occupies five percent of a brain region and one that occupies thirty percent is not just a numerical distinction; it can be the difference between watchful monitoring and emergency surgical intervention. This level of quantitative precision is what the system is designed to deliver, consistently and automatically, without requiring a specialist to manually trace boundaries or estimate proportions from a grayscale image. In doing so, it addresses one of the most persistent and practically consequential gaps in current CAD technology.

Deep learning, particularly through Convolutional Neural Networks, brings real analytical power to the classification side of this challenge. These models can identify patterns in MRI data that are genuinely imperceptible to the human eye — subtle textural differences, intensity gradients, and structural asymmetries that distinguish gliomas from meningiomas from pituitary tumors with a consistency and reliability that manual inspection cannot match at scale. But a diagnosis alone, however accurate, doesn't help a patient navigate the frightening and often confusing landscape of what comes next — what to eat, how active they can be, when to return for their next scan, and which specialist they need to see. That's why this system goes further than classification, translating its findings into practical, personalized care recommendations covering diet, physical activity, and follow-up scheduling, all calibrated to the specific tumor type identified and how advanced it appears to be. This extension from technical output to patient-centered guidance reflects a conviction that the value of a diagnostic tool should be measured not just by how accurately it classifies images but by how meaningfully it improves the experience and outcomes of the people it serves. A patient who leaves a diagnostic encounter with a clear understanding of their condition and a concrete plan for managing it is in a fundamentally better position than one who receives a label and a referral with no further context. The bigger picture here is about reshaping how

neuro-oncology works — by handling the time-consuming analytical groundwork, the system gives doctors back the mental bandwidth to focus on what only humans can truly do: making nuanced, empathetic decisions and building the kind of trust with their patients that no algorithm can replicate.

1.2 Problem Statement

Right now, diagnosing a brain tumor depends heavily on a radiologist manually reviewing MRI scans — a process that is slow, inconsistent, and offers no real-time diagnostic support at the point where it is often most urgently needed. This dependency on expert manual interpretation creates a fragile system where the quality and speed of diagnosis varies enormously based on who is available, how rested they are, and what imaging resources the facility has access to. In a field where early detection directly correlates with survival probability, this variability is not merely an operational inefficiency — it is a patient safety issue. Most existing tools only address part of this problem; very few bring together segmentation, classification, severity assessment, and patient guidance under one coherent roof, which means clinicians are often left to stitch together insights from multiple disconnected sources before they can form a complete clinical picture. The gaps this fragmentation creates are serious, and they show up consistently across five dimensions that this investigation has identified as the most clinically consequential. The first and most foundational gap is in tumor size estimation and segmentation. Most available tools cannot reliably delineate where a tumor ends and healthy tissue begins, particularly at the irregular, diffuse margins that are characteristic of malignant brain tumors, let alone calculate what proportion of the affected brain region the tumor occupies. This information is not incidental — it is central to staging a patient's condition and determining the appropriate intensity of treatment. Without it, clinicians are left estimating severity visually, introducing the kind of subjective variability that a quantitative system should eliminate. The second gap is in tumor classification: there is no widely available platform that reliably distinguishes between gliomas, meningiomas, and pituitary tumors in an automated, clinically deployable way, which means that treatment decisions — which differ substantially across these categories — continue to depend heavily on specialist availability and interpretation. The third gap is in automated reporting; generating a clear, structured, and complete diagnostic report remains largely a manual task in most clinical

settings, consuming specialist time, introducing inconsistency, and adding delays to a process where speed is often critical. The fourth gap is perhaps the most human one: after receiving a diagnosis, patients are frequently left without clear direction about what it means for their daily life. Most systems offer nothing in the way of lifestyle, dietary, or follow-up guidance tailored to the patient's specific condition, leaving people to navigate a bewildering post-diagnosis landscape largely on their own. The fifth gap is accessibility — many existing models are computationally demanding and too slow for practical use in real clinical workflows, and almost entirely unsuitable for the telemedicine platforms and under-resourced settings where fast, reliable automated diagnosis would arguably have the greatest impact on patient outcomes.

1.3 Objectives

The core goal of this project is to build an intelligent, deep learning-powered system that can detect, segment, and classify brain tumors from MRI scans — and then translate those findings into practical, personalized guidance that serves both patients and the healthcare professionals caring for them. This is a deliberately ambitious objective, because the problem it addresses is not a narrow technical one but a systemic clinical challenge that spans image analysis, decision support, and patient communication. Each of the system's five core objectives has been defined in direct response to one of the gaps identified in the problem statement, ensuring that the development effort is oriented toward real clinical need rather than abstract technical performance.

The first objective is accurate segmentation and sizing — using advanced preprocessing combined with fuzzy thresholding to cleanly isolate tumor regions from surrounding healthy tissue and express their extent as a meaningful percentage of the affected brain area. Achieving this requires handling the ambiguous, gradient-heavy boundaries that binary thresholding methods cannot manage reliably, and the fuzzy approach has been chosen specifically because it is mathematically suited to reasoning about exactly this kind of uncertainty. The second objective is reliable multi-class classification, training a CNN to accurately distinguish between glioma, meningioma, and pituitary tumors so that clinicians have a specific, actionable diagnosis to work from rather than a generic tumor-present alert. The third objective is clear, automated diagnostic reporting —

generating structured reports that cover tumor type, estimated size, severity stage, and recommended next steps in a format that is complete enough for clinical use and clear enough for a patient to understand without a medical degree. The fourth objective is personalized recommendations: offering concrete, condition-specific guidance on diet, daily activity levels, and follow-up consultation frequency that gives patients a practical roadmap for managing their health after diagnosis. The fifth objective is real-time usability — ensuring the system is lightweight and responsive enough to function in busy clinical environments, on mobile platforms, and in telemedicine contexts where computational resources may be limited and where every second of unnecessary delay has a human cost.

1.4 Scope of the Project

This system is designed to function as a complete, end-to-end diagnostic companion — one that doesn't just identify a problem and stop there, but actively helps map out what to do about it and supports both the clinical team and the patient through the steps that follow. The scope of the project has been defined to cover the full diagnostic journey from raw image input to personalized patient output, with each component of the pipeline designed to add value that builds on what came before it. Rather than optimizing any single capability in isolation, the system's scope reflects a deliberate commitment to integration — the conviction that a collection of individually strong components working seamlessly together is more clinically valuable than any one of them operating alone.

Segmentation and size estimation form the first pillar of the project scope. Fuzzy thresholding is applied to accurately isolate tumor regions from the surrounding brain tissue, and the resulting segmentation is used to calculate tumor size as a proportion of the total brain area — giving clinicians concrete, reproducible numbers rather than visual estimates that vary with the interpreter. Tumor classification forms the second pillar, with CNN-based models trained to sort detected tumors into the three clinically meaningful categories of glioma, meningioma, and pituitary tumor, ensuring that the diagnosis produced is specific enough to directly inform treatment planning. Automated diagnostic reporting represents the third pillar, with the system generating comprehensive reports that bring together tumor type, size percentage, severity stage, and recommended next steps into a single structured document that reduces manual workload and accelerates the path from analysis to clinical decision.

The fourth pillar extends the system's reach beyond the scan itself, providing customized lifestyle and health recommendations — covering diet, daily routines, activity levels, and follow-up scheduling — that support the patient's long-term wellbeing rather than treating the diagnostic report as the end of the interaction. The fifth and final pillar of the scope is scalable, real-time performance: the system has been architected to operate efficiently across clinical workstations, mobile applications, and telemedicine platforms, without imposing computational demands that would make it inaccessible to the smaller facilities and under-resourced settings that arguably have the most to gain from reliable automated diagnostic support.

1.5 Significance of the Study

What makes this project genuinely meaningful isn't the strength of any individual component — it's the way all of those components work together as a unified system, each one reinforcing the others in a way that produces clinical value greater than the sum of its parts. There are many tools in existence that can classify an MRI scan, and many others that can generate a report or offer health advice in isolation. What is far rarer — and what this investigation has set out to demonstrate is both achievable and practically deployable — is a platform that does all of these things coherently, in real time, in a single workflow that a clinician or patient can navigate without technical expertise. That integration is not just a design convenience; it has direct consequences for how efficiently care is delivered and how well-informed both doctors and patients are at every stage of the diagnostic and treatment process.

The significance begins with the simple but transformative idea of a unified diagnostic platform. Doctors working in busy clinical environments currently navigate a fragmented landscape of tools, each addressing a different piece of the diagnostic puzzle, with the clinician responsible for manually connecting the dots between them. This system replaces that fragmentation with a single, coherent platform where everything from tumor size to recommended follow-up intervals is available in one place, reducing cognitive load and the risk of important information falling through the gaps. Better detection and staging is the second dimension of significance — by calculating tumor size as an objective, reproducible proportion of brain tissue and categorizing it by type and severity, the system gives clinicians a far richer quantitative foundation for treatment planning than visual interpretation alone can provide. Smarter, automated reporting addresses the third dimension:

detailed, consistently structured reports reduce the risk of diagnostic oversight, accelerate the speed at which clinical decisions can be made, and — perhaps most importantly for the patient — communicate findings in language that is clear and accessible rather than locked behind specialist terminology. Holistic patient care represents the fourth dimension of significance, reflecting the conviction that a diagnostic tool's responsibility doesn't end at the classification result. By pairing its findings with tailored dietary advice, lifestyle guidance, and structured follow-up scheduling, the system acknowledges that the patient is a whole person navigating a frightening situation, not merely a scan to be analyzed and filed. The fifth and final dimension of significance is the strengthening of clinical collaboration — clear, actionable outputs that bridge the communication gap between what the AI has found and what the patient needs to understand create the conditions for earlier intervention, more informed shared decision-making, and ultimately better outcomes for the people at the center of it all.

CHAPTER 2
LITERATURE SURVEY

2.1 Literature Survey

1. Sasikala and Kumaravel (2023) – "Wavelet-Based Texture Features and Genetic Algorithm with Neural Networks for Brain Tumor Classification," International Conference on Medical Imaging, India, 2023, pp. 45-50.

This research took a fresh approach to brain tumor classification by combining wavelet-based texture analysis with Genetic Algorithms for feature selection, all tied together through a neural network for the final predictions. The core idea was straightforward: rather than throwing every piece of data at the model, the researchers wanted to work smarter by identifying only the most meaningful features before classification.Starting with pre-processed MRI scans, the team applied wavelet transforms to pick up on subtle textural patterns in brain tissue — details that might easily go unnoticed with the naked eye. From there, Genetic Algorithms helped trim the fat, weeding out redundant information and keeping only what truly mattered for accurate tumor detection. Those refined features were then handed off to a neural network to make the final call.The results were encouraging. The model outperformed traditional statistical methods in terms of accuracy, which was a meaningful step forward. That said, the research wasn't without its shortcomings.When it came to distinguishing between multiple tumor types — glioma, meningioma, and pituitary tumors — the model struggled to keep up. Real-time processing was another missing piece, which matters enormously in clinical settings where speed can be critical.

2. El-Dahshan et al. (2023) – "Hybrid Approach Combining DWT, PCA, and Neural Networks for Brain Tumor Diagnosis," Journal of Computational Intelligence in Medicine, 2023, pp. 33-39.

El-Dahshan and his team tackled brain tumor detection by bringing together three well-established techniques into a single hybrid system. Their pipeline combined Discrete Wavelet Transform (DWT) for breaking down MRI images, Principal

Component Analysis (PCA) for keeping the feature set manageable, and classifiers — KNN and Feedforward Neural Networks — to handle the actual tumor identification. The thinking behind it was sound: by layering these methods together, they hoped to get more accurate results than any single approach could deliver on its own.

The process worked like this — MRI scans were first decomposed using DWT to capture information from both the spatial and frequency domains. PCA then stepped in to cut down the number of features, reducing computational load without sacrificing too much detail. From there, the classifiers took over and made the final predictions.

On accuracy, the system delivered. It measurably outperformed older, more conventional image processing methods, which was the primary goal.But the cracks started to show when looking beyond raw accuracy. The multiple pre-processing stages piled up computationally, making the system too slow for time-sensitive clinical environments. There was also no segmentation built in, meaning the system could flag a tumor's presence but couldn't map its boundaries or estimate its size — critical information for any treating physician.

3. Kaplan et al. (2024) – "LBP and Texture-Based Feature Extraction for Brain Tumor Analysis," International Symposium on Biomedical Imaging, 2024, pp. 120-126.

Kaplan and his team took a texture-first approach to brain tumor classification, building their method around a family of Local Binary Pattern descriptors — LBP, nLBP, and αLBP. The idea was to let these texture-based features do the heavy lifting when it came to capturing meaningful patterns in MRI images, which were then fed into a trio of classifiers: KNN, Artificial Neural Networks, and Random Forest. The result was an accuracy of 95.56%, a figure that comfortably beat out traditional statistical models and gave the team good reason to be confident in their approach.What the study made clear is that texture descriptors can be genuinely powerful tools, particularly when the task involves telling apart healthy brain tissue from abnormal growths. In that specific context,

The entire pipeline rests on handcrafted features, which means someone has to carefully engineer what the system looks for — there's no automatic learning

happening under the hood. This is precisely where modern deep learning models have pulled ahead, learning rich representations directly from raw data without needing that manual groundwork. As MRI quality and imaging protocols vary from one hospital to the next, a system so dependent on fixed feature engineering starts to struggle with adaptability.Then there's the clinical gap. Without any segmentation capability, the system can't draw tumor boundaries or estimate size — two things that a clinician genuinely needs before making any treatment decision. There's no severity analysis either, which further limits how useful the output actually is in a medical setting.

.

4. Mohsen et al. (2024) – "Deep Neural Networks for Brain Tumor Detection Using PCA and DWT," Journal of AI in Healthcare, 2024, pp. 89-94.

Mohsen and his team set out to bridge the gap between traditional image processing and deep learning by combining PCA and DWT for feature extraction with a Deep Neural Network for the final classification. The underlying ambition was clear — take the structured, well-understood power of conventional feature extraction and pair it with the representational strength of a DNN to get the best of both worlds.

And to a degree, it worked. The model showed genuine promise, demonstrating that thoughtfully integrated deep learning approaches could outshine shallower machine learning methods. That in itself was a meaningful contribution at the time.

But the study's limitations are difficult to set aside, and some of them are fairly fundamental.The most glaring issue is the dataset — just 66 MRI images. That's a remarkably small foundation on which to build and evaluate any model, let alone one intended for medical use. With so little data, overfitting becomes almost inevitable, and any accuracy figures need to be read with that in mind. A model trained on 66 images simply hasn't seen enough variation to be trusted in a clinical environment where no two patients are alike.

5. Cheng (2023) – "Support Vector Machines with Texture Features for Brain Tumor Classification," Neurocomputing Advances, 2023, pp. 140-145.

Cheng made a notable contribution to the field by introducing a dedicated dataset for brain tumor classification — something that in itself was valuable for the research

community. Alongside this, he proposed a classification approach built around Support Vector Machines, enriched with handcrafted features drawn from histograms and Gray Level Co-occurrence Matrix (GLCM) analysis. The system was designed with texture-based patterns at its core, targeting three distinct tumor categories. Hitting an accuracy of 91.28%, the model proved that well-crafted statistical features paired with a solid machine learning backbone could deliver respectable results, particularly when working with smaller datasets.The study made a convincing case for the continued relevance of classical feature extraction methods at a time when the field was still finding its footing with more advanced approaches.

There was also no automated reporting or recommendation mechanism, which makes integration into a real clinical workflow difficult to imagine.Cheng's work remains a meaningful piece of early-stage research and a useful dataset contribution, but measured against the demands of today's AI-driven healthcare landscape, it highlights just how far the field has come — and how much further it still needed to go.

6. Pashaei et al. (2024) – "Hybrid CNN and Kernel Extreme Learning Machines for Brain Tumor Detection," IEEE Transactions on Neural Systems, 2024, pp. 178-183.

This study brought together the feature-learning strength of Convolutional Neural Networks with the speed of Kernel Extreme Learning Machines to create a hybrid classification system. The pairing made practical sense — CNNs handled the heavy lifting of extracting meaningful patterns from MRI images, while KELM kept the classification process fast and efficient. The outcome was a solid 91.28% accuracy with noticeably reduced computation time compared to more traditional approaches, and a meaningful step away from the manual feature engineering that had held earlier methods back.

That said, the model's shortcomings are worth noting.Tumor segmentation and size estimation were both absent, which makes the system difficult to lean on for actual treatment planning. The dataset used was relatively limited in diversity, raising genuine questions about how well the model would hold up across different hospitals and imaging conditions. Clinical interpretability was another gap — without

explainable outputs, gaining the trust of medical professionals becomes an uphill battle.

There was also no provision for automated reporting or health recommendations, keeping the system firmly in the classification-only lane. Incorporating multi-task learning to handle both segmentation and classification together could be a strong direction for future work, pushing this kind of hybrid approach closer to something clinically meaningful.

7. Phaye et al. (2025) – "Capsule Networks for Brain Tumor Classification," ACM Conference on Medical Image Computing, 2025, pp. 210-215.

Phaye and his team used Capsule Networks (CapsNet) to classify brain tumors from MRI images, with the model achieving an impressive 95.03% accuracy. Unlike traditional CNNs, CapsNet excels at preserving spatial relationships, which helped it better distinguish between tumor types like glioma, meningioma, and pituitary tumors.

However, the model came with notable drawbacks. Training was slow and computationally demanding, making it impractical for real-time clinical use. It also lacked key features like tumor segmentation, size estimation, and automated reporting — all of which matter greatly in actual treatment planning. While CapsNet shows real promise for future AI-driven medical tools, its high resource requirements remain a significant hurdle, especially in healthcare settings with limited infrastructure.

8. Narasimham et al. (2023) – "Fuzzy Logic-Based Binary Classification of Brain MRI," International Journal of Computational Imaging, 2023, pp. 77-81.

This study used fuzzy logic to classify MRI scans as either tumor or non-tumor, reaching a solid 96% accuracy. Its biggest strength was simplicity — the decision-making process was easy to follow and lightweight enough for faster inference without heavy computational demands. It also handled uncertainty in medical imaging reasonably well, which is a common challenge in this field. That said, the model had clear limitations. It could only perform binary classification, meaning it

had no ability to distinguish between tumor types like glioma or meningioma. The dataset used was also very small — just 253 images — which raises questions about how well it would perform on broader, real-world data. On top of that, it lacked segmentation, tumor sizing, deep learning integration, and automated reporting, all of which are increasingly expected in modern clinical AI tools. Ultimately, while fuzzy logic offers interpretability and ease of use, it falls short for anything beyond basic screening tasks.

9. Van Hai and Amaechi (2024) – "Fuzzy Logic-Based Segmentation Combined with CNN for Brain Tumor Detection," IEEE International Conference on Image Processing, 2024, pp. 245-250.

Van Hai and Amaechi developed a hybrid model that combined fuzzy logic-based segmentation with CNN-based classification, achieving an accuracy of 97.6%. By first segmenting the tumor region and then classifying it, the system added real clinical value — particularly through its ability to visualize tumor boundaries, which aids in diagnosis.

However, the model still had meaningful gaps. It lacked tumor size estimation and severity analysis, both of which are critical when planning treatment. The dataset was small and not diverse enough to confidently support real-world performance. Missing features like automated reporting and personalized patient recommendations also held it back from being a complete clinical tool. While the combination of segmentation and classification was a genuine step forward, the system fell short of the end-to-end decision support that modern hospitals and telemedicine platforms truly need.

10. Maqsood et al. (2025) – "Fuzzy Logic and U-Net CNN for Brain Tumor Segmentation," Neural Computing and Applications Journal, 2025, pp. 310-315.

Maqsood and colleagues merged fuzzy logic with the U-Net CNN architecture to tackle brain tumor segmentation, reaching an impressive 98% efficiency rate. The combination worked well — U-Net handled pixel-level detail while fuzzy logic managed adaptive thresholding, together minimizing information loss along tumor

boundaries.That said, the system was built solely for segmentation and never extended into classification. There was no way to identify whether a tumor was glioma, meningioma, or any other type. It also lacked tumor size estimation, clinical reporting, and any form of decision-support functionality, which are all essential for real diagnostic use. The limited dataset size further raised concerns about how reliably the model would perform across varied MRI scans in practice. While the U-Net and fuzzy logic pairing proves to be a strong foundation for segmentation, future work clearly needs to go further — integrating classification, automated reporting, and clinical recommendations to build a truly comprehensive brain tumor diagnosis system.

2.2 Inferences and Challenges in Existing Systems

The existing systems for brain tumor detection have long leaned on radiologists manually combing through MRI scans and applying traditional image processing techniques — a workflow that, while functional, comes with serious limitations in accuracy, speed, and the ability to scale across busy healthcare environments. Most current approaches rely on relatively straightforward methods like threshold-based segmentation or texture feature extraction, paired with classical classifiers such as SVM, KNN, or ANN. These techniques aren't without merit — they've served the field for years — but they were never designed to handle the complexity and variability that real-world brain tumor cases demand. One of their most glaring shortcomings is the inability to precisely map tumor boundaries or estimate volume with the kind of reliability that treatment planning requires. The situation becomes even more limiting when you consider that the majority of these models are built only for binary classification — essentially answering "is there a tumor or not?" — without any ability to distinguish between glioma, meningioma, or pituitary tumor types. That distinction isn't a minor detail; it fundamentally shapes what treatment a patient receives and how urgently they need it. Without integrated size calculation or severity grading, clinicians are left to fill in the blanks themselves, adding cognitive load and potential for error to an already demanding process.

Inferences from Literature

A close look at the existing body of research on brain tumor detection reveals a field that has made genuine progress but is still working through some fundamental

limitations. Traditional methods that rely on manual interpretation or classical machine learning remain the norm in many settings, and while experienced radiologists bring irreplaceable clinical intuition to the table, these approaches are inherently time-consuming and vulnerable to fatigue-related inconsistencies. Segmentation techniques like thresholding and wavelet-based methods have shown reasonable performance under controlled conditions, but they tend to struggle when tumors present irregularly — which, in practice, is more often than not. Fuzzy logic and deep learning approaches have emerged as stronger alternatives precisely because they can handle this kind of ambiguity more gracefully, treating the uncertain, gradient-heavy regions at tumor margins as mathematically manageable rather than simply problematic. CNN-based models, in particular, have demonstrated a clear edge over shallower architectures in both accuracy and feature extraction depth, suggesting that deep learning isn't just a trend in this space — it's becoming a baseline expectation for serious diagnostic tools. Yet even many deep learning systems fall into the same trap as their predecessors: binary classification that tells you a tumor exists but nothing about what kind it is, how large it is, or what stage it may have reached.

The research also consistently reveals a striking absence of severity grading and tumor size estimation — two pieces of information that are arguably as important as detection itself when it comes to staging a patient's condition and planning treatment. Integrated reporting mechanisms that combine clinical findings with patient-friendly summaries and actionable next steps are almost entirely absent from the literature, which means patients often walk away from a diagnosis without any real understanding of what it means for their day-to-day life. Explainability is another recurring gap; most models operate as black boxes, and the medical community's reluctance to adopt tools it can't interrogate is both understandable and well-founded. Real-time performance and lightweight deployment are similarly overlooked, leaving many technically impressive solutions confined to research settings rather than functioning in the hospitals and telemedicine platforms that actually need them. The lack of personalized health guidance — diet modifications, lifestyle adjustments, structured follow-up scheduling — means these systems treat diagnosis as the finish line when, for the patient, it's really just the beginning. Finally, scalability and cross-institutional adaptability remain persistent challenges; a model that performs well on one hospital's imaging dataset often struggles to generalize

when confronted with the different equipment, protocols, and patient demographics of another. Taken together, these inferences paint a clear picture of where the field needs to go: toward integrated, explainable, patient-centered systems that work reliably across diverse real-world conditions.

Challenges in Existing Systems

Building effective brain tumor detection systems isn't just a technical challenge — it's a deeply human one, and the obstacles that current systems face reflect just how complex the intersection of medicine, data science, and clinical practice really is. One of the most immediate technical hurdles is precise segmentation: delineating exactly where a tumor ends and healthy brain tissue begins is genuinely difficult, especially when tumor margins are diffuse or when surrounding edema creates ambiguity in the imaging. Without accurate boundary detection, any downstream analysis — size estimation, severity grading, treatment planning — is built on an uncertain foundation, which has direct consequences for patient safety. Most current systems compound this problem by stopping at detection without offering any contextual information about the tumor's size relative to the brain or what stage it may represent, leaving clinicians to source that critical information elsewhere. Deep learning models, which hold the most promise for overcoming these limitations, introduce their own complications: they require large, carefully annotated datasets to train effectively, and in the domain of brain MRI, such datasets are both scarce and frequently imbalanced across tumor types. A model trained predominantly on glioma cases, for instance, may perform poorly when presented with a meningioma, creating reliability gaps that are unacceptable in a clinical context. Interoperability adds another layer of difficulty — most research prototypes exist in isolation and have no straightforward way of connecting with the hospital information systems and Picture Archiving and Communication Systems (PACS) that clinicians actually use every day.

The computational demands of high-performing models present a particularly thorny challenge for real-world deployment. Most state-of-the-art architectures require significant processing power that simply isn't available in smaller clinics or community health centers, and this creates a troubling disparity where the most advanced diagnostic tools are accessible only in well-resourced institutions. Explainability remains a critical sticking point — medical professionals need to

understand why a system reached a particular conclusion before they can responsibly act on it, and the opacity of most current AI models makes this nearly impossible. The absence of personalized post-diagnosis recommendations means patients are often left without a clear roadmap after receiving their results, unsure of what to eat, what to avoid, how active they can be, or when to return for their next scan. Variability in MRI quality and acquisition protocols across different hospitals introduces further complexity, as models trained in one imaging environment can degrade significantly when exposed to scans captured under different conditions. Data privacy and security concerns layer on top of all of this, since patient MRI scans and associated medical histories are among the most sensitive categories of personal data and must be handled in strict compliance with healthcare regulations like HIPAA and GDPR. Finally, almost none of the existing systems have been designed with large-scale deployment in mind — they work well enough as research demonstrations but haven't been engineered for the demands of a major hospital network or a national telemedicine platform, which means the gap between what's technically possible and what's practically available to patients remains frustratingly wide.

CHAPTER 3
AIM AND SCOPE OF THE PRESENT INVESTIGATION

3.1 Aim of the Study

The primary aim of this investigation is to design and build a complete, end-to-end automated framework for detecting, classifying, and clinically assessing brain tumors using advanced Computer Vision and Deep Learning techniques. Medicine is rapidly moving toward digital health solutions, and with that shift comes an urgent need for diagnostic tools that are fast, non-invasive, and genuinely reliable under real clinical pressure. Brain tumor diagnosis has historically been a slow and labor-intensive process — one where the gap between a patient receiving an MRI scan and actually beginning treatment can stretch into days or even weeks, a delay that in oncology can carry serious consequences. This project is fundamentally about compressing that gap. By providing an immediate, AI-driven secondary opinion to radiologists, the system aims to give clinicians a head start — something they can cross-reference their own interpretation against before committing to a diagnostic conclusion. The goal isn't to replace the radiologist but to function as a tireless, highly consistent analytical partner that catches what fatigue or visual limitation might cause a human to miss. Beyond the clinical setting, this kind of tool has profound implications for telemedicine and under-resourced healthcare environments where specialist radiologists simply aren't available around the clock. Ultimately, this investigation is driven by a straightforward but ambitious conviction: that the right combination of deep learning, image processing, and intelligent reporting can meaningfully shorten the distance between a scan and a life-saving intervention.

3.1.1 Precision Multi-Class Classification

One of the most important aims of this study is to push past the limitations of simple binary tumor detection and build something far more clinically useful. Most existing systems answer only one question — is there a tumor present or not — and while that's a meaningful starting point, it leaves physicians without the specific information they need to plan a course of treatment. This investigation focuses on classifying brain MRI scans across four distinct categories: Glioma, Meningioma,

Pituitary tumors, and healthy brain tissue with no tumor present. Each of these conditions has a different biological origin, a different pattern of growth, and a very different set of treatment options, which means distinguishing between them isn't just academically interesting — it's clinically essential. To achieve this, a Deep Convolutional Neural Network has been trained specifically on the structural and intensity characteristics that differentiate these tumor types at the pixel level. These differences can be extraordinarily subtle; a glioma and a meningioma might look broadly similar to an untrained eye, but a well-trained CNN can learn to pick up on the textural, boundary, and contrast cues that separate them. The network is designed to automatically extract these hierarchical features layer by layer, building from basic edge detection in the early layers to highly abstract tumor-specific representations deeper in the architecture. This approach mirrors, in a computational sense, the way an experienced radiologist develops pattern recognition over years of reading thousands of scans — except the CNN can apply that learned pattern recognition instantly and consistently, without variation due to fatigue or workload. The result is a classification system capable of delivering a specific, confident diagnosis rather than a vague alert, giving both physicians and patients a much clearer picture of what they're dealing with from the very first analysis.

3.1.2 Integration of Fuzzy Logic for Segmentation

A particularly distinctive aim of this investigation is the incorporation of Fuzzy Thresholding into the image segmentation process, setting it apart from the majority of existing detection systems that rely on more rigid, binary approaches. Traditional global thresholding works by drawing a hard line — pixels above a certain intensity value are classified as tumor, and everything below is classified as healthy tissue. In theory, this sounds straightforward. In practice, it fails regularly because the boundary between tumor and healthy tissue is rarely clean. Tumors are biological structures, and their edges often blend gradually into surrounding edema or inflamed tissue, creating intensity gradients that a hard threshold simply cannot handle with precision. Fuzzy logic is designed exactly for this kind of ambiguity. Rather than forcing every pixel into one of two categories, fuzzy thresholding assigns each pixel a membership value — a measure of how likely it is to belong to the tumor region — and uses that probability to make more nuanced segmentation decisions. This

allows the system to capture the full extent of a tumor, including its diffuse outer margins, without over-segmenting into healthy tissue or under-segmenting and missing critical portions of the growth. The clinical significance of this precision is hard to overstate; a surgeon planning a resection needs to know exactly how much tissue is involved, and even a small segmentation error can translate into a poorly planned procedure. By demonstrating that fuzzy logic outperforms global thresholding in handling these uncertain boundary regions, the study aims to make a meaningful methodological contribution to the broader field of medical image analysis. The goal, ultimately, is segmentation quality that approaches surgical-grade accuracy — detailed enough to be genuinely useful in a pre-operative planning context rather than just good enough for a research paper.

3.1.3 Intelligent Staging and Severity Assessment

Beyond detecting and classifying a tumor, this investigation aims to introduce a module that brings a quantitative dimension to the diagnostic picture — something that existing systems almost universally overlook. Once the tumor region has been segmented from the MRI scan, the system calculates the pixel area of that region and expresses it as a proportion of the total brain area visible in the scan. This produces a concrete, numerical estimate of how much space the tumor is occupying, which the system then uses to automatically categorize the tumor into one of three severity stages: Small, Medium, or Large. While these categories are deliberately straightforward, the information they encode is anything but trivial. An oncologist deciding whether a patient needs urgent surgical intervention, a course of radiotherapy, or a watchful waiting approach needs some measure of tumor burden to work from, and right now that information is either estimated visually or derived from time-consuming manual measurements. Automating this step doesn't just save time — it standardizes the measurement, removing the variability that comes from different clinicians applying different mental benchmarks for what constitutes a "large" tumor. The staging output also feeds directly into the recommendation engine, ensuring that the lifestyle and follow-up guidance the system generates is appropriately calibrated to the severity of the patient's condition. A patient with a small, slow-growing pituitary tumor needs very different follow-up scheduling than someone with a large, rapidly expanding glioma, and the system is designed to reflect that difference.

3.1.4 Automated Reporting and Personalized Recommendations

One of the most human-centered aims of this project is to automate the reporting process and extend the system's output beyond clinical findings into genuinely practical patient guidance. In most diagnostic workflows, generating a written report is a time-consuming task that falls entirely on the radiologist or attending physician, adding to an already heavy administrative burden and introducing delays between analysis and communication. This system changes that by automatically generating a structured clinical report the moment image processing is complete, capturing the tumor type, size estimate, severity stage, and relevant observations in a format that is both medically precise and readable by a non-specialist audience. The report is designed to serve two audiences simultaneously — the clinician, who needs accurate technical data to make treatment decisions, and the patient, who deserves to understand their own condition in language that doesn't require a medical degree to parse. Beyond the report itself, the system generates a set of personalized recommendations tailored to the specific tumor type and severity identified. On the dietary side, this means providing nutrient-rich food guidance appropriate for patients experiencing neurological stress — foods that support brain health, reduce inflammation, and complement the metabolic demands of treatment. The lifestyle guidance component addresses physical activity levels, sleep habits, and environmental adjustments that can meaningfully affect a patient's quality of life and recovery trajectory during and after treatment. The system also includes medical consultation routing, suggesting the appropriate specialist — whether a neurosurgeon, radiation oncologist, or neurologist — based on what the scan has revealed, so that patients aren't left unsure of which door to knock on next. Taken together, these features represent a shift in how AI diagnostic tools are conceived: not as systems that produce a finding and then step back, but as platforms that stay engaged with the patient's care journey from the moment of diagnosis onward. The ultimate ambition here is to reduce the anxiety and confusion that so often follow a brain tumor diagnosis by replacing uncertainty with clear, actionable information that both doctors and patients can act on immediately.

3.1.5 Real-Time Accessibility and Deployment

A diagnostic tool is only as valuable as its ability to reach the people who need it, and one of the most practical aims of this investigation is to ensure that the system

can run effectively without requiring expensive, specialized hardware. Many of the most powerful deep learning models in medical imaging research exist only in high-performance computing environments — GPU clusters, cloud servers, or dedicated workstations that cost far more than the average clinic in a mid-sized or rural city can reasonably budget for. This creates a troubling divide where cutting-edge AI capabilities are concentrated in wealthy, well-equipped institutions while smaller facilities and telemedicine platforms — often the ones serving the most underserved populations — are left behind. This project directly addresses that gap by optimizing the CNN architecture for efficiency and deploying the system through the Gradio framework, a lightweight, browser-based interface that can run on a standard laptop or tablet without any specialized setup. The interface is deliberately simple and intuitive, designed so that a clinician with no background in machine learning can upload a scan, receive a full classification and report, and navigate the output without any technical friction. Response times are optimized so that predictions are delivered within seconds of image submission, maintaining the kind of pace that a real clinical workflow actually demands. The Gradio deployment also means the system can be accessed remotely, making it well-suited for telemedicine scenarios where a radiologist might be reviewing scans from a different city or even a different country. By proving that high-quality AI diagnostics don't have to be locked behind expensive infrastructure, this investigation aims to make a practical argument for broader adoption — demonstrating that the technology is not only powerful but genuinely deployable in the messy, resource-constrained reality of everyday healthcare.

3.2 Scope of the Study

The scope of this investigation covers the entire diagnostic pipeline from the moment a raw MRI image enters the system to the point at which a complete clinical report is delivered to the end user. Rather than focusing on one narrow slice of that pipeline — as many research projects do — this study is designed to evaluate how well all the components work together as an integrated whole. The investigation is specifically grounded in the computational challenges of neuro-oncology, working exclusively with T1-weighted and T2-weighted MRI scans, which remain the standard imaging modalities for brain tumor assessment in clinical practice. Every design decision made in the study — from architecture choices to preprocessing

strategies to interface design — is oriented toward the practical realities of clinical deployment rather than abstract benchmark performance. The scope is intentionally ambitious but clearly bounded, ensuring that the research produces something complete and evaluable within its defined domain rather than partially addressing a much larger problem. This chapter outlines precisely what the investigation covers, what methods it employs, and where it draws its boundaries, so that the findings can be understood and built upon by future researchers with full clarity about what has and hasn't been addressed.

3.2.1 Data Domain and Limitations

The investigative domain of this study is strictly confined to brain MRI images, and this boundary has been drawn deliberately rather than arbitrarily. CT scans, X-rays, and PET scans each carry different imaging characteristics and would require entirely separate preprocessing pipelines and model architectures to analyze effectively — including them would dilute the study's focus without meaningfully improving its outputs. The dataset used in this investigation comprises over 7,000 high-resolution MRI images, a volume large enough to provide statistically significant training signal and to give the model meaningful exposure to the natural variation in how tumors present across different patients, ages, and scan conditions. The four classes covered — Glioma, Meningioma, Pituitary tumor, and No Tumor — represent the most commonly encountered categories in clinical brain tumor diagnosis and collectively account for the vast majority of cases a neurological imaging department would encounter in practice. That said, the investigation does not extend to rare metastatic brain diseases, neuro-degenerative conditions like Alzheimer's disease, or other pathologies that fall outside these four categories. This is a deliberate scope limitation rather than an oversight; incorporating rare conditions without sufficient representative data would risk training an overconfident model that performs poorly on precisely the cases where diagnostic mistakes are most costly. The dataset has also been curated to maintain reasonable class balance, though the natural prevalence differences between tumor types introduce some inherent imbalance that the study addresses through augmentation strategies discussed in later sections.

3.2.2 Algorithmic Scope

The technical heart of this investigation is the design and training of a custom Convolutional Neural Network architecture built specifically for the demands of multi-class brain tumor classification. Rather than relying on a pre-built, off-the-shelf architecture, the network has been designed from the ground up to balance classification accuracy with the computational efficiency needed for real-time deployment on standard hardware. The feature extraction process works through a series of convolutional layers that learn spatial hierarchies automatically from the raw pixel data — early layers pick up on basic visual features like edges and intensity gradients, while deeper layers combine these into increasingly complex and tumor-specific representations. This automatic feature learning is one of the fundamental advantages CNNs hold over classical machine learning approaches, which require domain experts to manually engineer the features the model will use, a process that is both time-consuming and inherently limited by human intuition about what matters visually. Dimensionality reduction is handled through MaxPooling layers, which progressively compress the spatial information extracted by the convolutional layers, retaining the most informative features while reducing the computational load on subsequent layers. Dropout layers are incorporated at strategic points in the architecture to combat overfitting — a persistent risk when training on medical datasets that, despite their size, are small by the standards of general computer vision tasks. The model is trained using the Adam optimizer, chosen for its adaptive learning rate properties that make it particularly well-suited to the non-stationary loss landscapes typical of deep network training. Sparse Categorical Crossentropy serves as the loss function, handling the multi-class label structure cleanly without requiring one-hot encoding of the target labels. Together, these design choices reflect a coherent architectural philosophy: build a model that learns the right things, generalizes well to unseen data, and runs efficiently enough to be genuinely useful outside of a research lab.

3.2.3 Image Processing and Hybrid Techniques

The scope of this investigation extends well beyond the neural network itself to encompass a comprehensive image preprocessing pipeline that prepares raw MRI scans for analysis in a way that maximizes the model's ability to learn meaningful features. Raw MRI images straight from the scanner are rarely in a condition suitable

for direct input into a deep learning model — they vary in brightness, contrast, and scale in ways that would introduce noise into the training process if left unaddressed. Normalization is the first step in resolving this, scaling all pixel intensity values to a standardized range between 0 and 1, which accelerates model convergence and ensures that no single image dominates the training signal simply because it happens to be brighter or higher in contrast than the others. Following segmentation, morphological operations including erosion and dilation are applied to the segmented tumor region to clean up the output — erosion removes small, isolated noise artifacts at the boundary of the detected region, while dilation ensures that the tumor outline is continuous and complete rather than fragmented. This combination of operations refines the segmentation result in a way that makes it more clinically meaningful and visually interpretable. Data augmentation plays a critical role in expanding the effective size and diversity of the training dataset by generating modified versions of existing images through controlled transformations including rotations, horizontal and vertical flips, and zoom variations. These transformations simulate the kinds of differences that naturally arise in real-world MRI data — different patient positioning, different scanner orientations, slight differences in how slices are taken — and expose the model to a wider range of presentations during training without requiring additional labeled data. The result is a model that is considerably more robust to the variability it will inevitably encounter when deployed in a real clinical setting, where scan conditions are rarely as consistent as a curated research dataset. Together, these preprocessing and augmentation strategies form an essential scaffolding around the deep learning model, ensuring that what the network learns reflects genuine diagnostic signal rather than dataset-specific artifacts.

3.2.4 Deployment and Interaction

The scope of this investigation deliberately extends into the user experience layer, recognizing that a technically excellent model is of limited value if it's difficult to actually use in practice. The deployment platform chosen for this system is the Gradio library, a Python-based framework that enables the rapid development of clean, browser-accessible interfaces for machine learning models without requiring any web development expertise. The interface is designed around the simplest possible user journey: a clinician or technician drags and drops an MRI image into

the upload area, and within seconds receives a complete classification result alongside the processed, segmented image. This immediacy is not just a convenience feature — in a clinical context, every second of unnecessary waiting adds friction to an already pressured workflow, and a system that takes minutes to return a result is one that clinicians will stop using. The visualization component is particularly important from a trust perspective. Rather than presenting only a text label or a confidence percentage, the system displays the segmented tumor region overlaid on the original scan, giving the clinician a visual reference that allows them to evaluate the AI's conclusion against their own interpretation of the image. This transparency is what transforms the system from a black box into a collaborative tool — one that shows its reasoning rather than simply announcing a verdict. The interface has also been designed to be accessible on tablet devices, recognizing that clinical environments don't always offer a desktop workstation at the point of care. By building usability directly into the scope of the investigation rather than treating it as an afterthought, the study ensures that its findings are relevant not just to other researchers but to the actual healthcare professionals the system is built to serve.

3.2.5 Performance Evaluation Metrics

Ensuring that the system meets a standard rigorous enough for medical application requires an evaluation framework that goes considerably further than simply reporting overall accuracy, and the scope of this investigation reflects that requirement. Confusion matrices form the backbone of the performance evaluation, providing a detailed breakdown of how the model's predictions align with ground truth labels across all four classes. This visualization makes it immediately apparent where the model's errors cluster — whether, for instance, it is systematically confusing glioma with meningioma, or whether it is reliably identifying healthy brain tissue while struggling with pituitary tumors. These patterns are clinically meaningful; a false negative on a malignant glioma carries a very different risk profile than a false positive on a benign meningioma, and the confusion matrix makes these distinctions visible in a way that aggregate accuracy cannot. Accuracy and loss curves tracked across 50 or more training epochs provide a longitudinal view of how the model's learning stabilizes over time, revealing whether it is converging smoothly or exhibiting signs of overfitting or underfitting that would compromise its reliability on

new data. Classification reports provide the granular precision and recall metrics that are especially important in a medical context — precision measures how often a positive prediction is correct, while recall measures how reliably the system finds all instances of each tumor type in the dataset. In neuro-oncology, recall is often the more critical of the two: missing a tumor that is present is generally a far more serious failure than flagging a scan for further review when no tumor exists. By evaluating the system across all of these dimensions simultaneously, the investigation builds a multi-faceted picture of performance that is honest about both the system's strengths and the areas where further development would be most impactful.

3.2.6 Exclusions and Future Outlook

The boundaries of this investigation have been drawn with as much care as its inclusions, and understanding what the system does not do is just as important as understanding what it does. Real-time video stream analysis — the continuous processing of MRI data as it is being acquired — falls outside the current scope, as this would require specialized hardware integration and a real-time data pipeline that goes well beyond what the present project infrastructure can support. Similarly, 3D volumetric rendering from DICOM files, while a genuinely valuable capability for surgical planning, requires dedicated computational resources and a level of 3D image processing complexity that represents a distinct research undertaking in its own right rather than a natural extension of the current system. These exclusions are practical rather than philosophical — they represent genuine next steps for the technology rather than capabilities deemed unimportant. Perhaps the most important boundary the system maintains is an ethical one: all outputs generated by the platform are explicitly framed as recommendations and analytical support rather than prescriptions or definitive medical conclusions. The final clinical decision in every case remains with the licensed neurosurgeon or treating physician, and the system is designed to reinforce rather than circumvent that responsibility.

CHAPTER 4
EXPERIMENTAL MATERIALS, METHODS, AND ALGORITHMS USED

4.1 Dataset Description and Data Engineering

The foundational integrity of the brain tumor detection system is fundamentally tied to the quality, diversity, and volume of the underlying MRI dataset. For this research, a large-scale, multi-institutional dataset was curated, comprising 7,023 high-resolution T1-weighted contrast-enhanced Magnetic Resonance Imaging (MRI) scans. These scans are meticulously categorized into four primary diagnostic classes: glioma, meningioma, pituitary tumors, and a control group representing "no tumor." Gliomas represent aggressive primary brain tumors, meningiomas typically arise from the membranes surrounding the brain, and pituitary tumors affect the endocrine system; each presents distinct morphological features that the model must learn to distinguish. The inclusion of the "no tumor" class is critical for clinical utility, as it reduces false-positive rates by training the network on healthy brain anatomy.

The data acquisition phase involved standardizing diverse image formats into a unified structure suitable for deep learning. The data is partitioned into two primary subsets through a stratified sampling approach: the training set, comprising 5,712 images, and the testing set, comprising 1,311 images. This stratification ensures that the model is exposed to a representative distribution of tumor types and anatomical variations during the learning phase. Furthermore, data engineering techniques such as pixel intensity normalization and metadata cleaning are applied to ensure that the convolutional filters focus on pathological significance rather than hardware-induced noise or varying brightness scales between different MRI scanners. By utilizing such a significant sample size, the system achieves the robustness necessary to handle variations in patient age and tumor stage, allowing the model to capture intricate spatial relationships that are often subtle to the human eye.

4.2 Advanced Image Preprocessing and Fuzzy-Logic Segmentation

Raw MRI data often contains artifacts, varying intensity scales, and electronic noise

that can impede the performance of convolutional filters. Therefore, a comprehensive preprocessing pipeline is established to ensure data consistency across all samples. All raw MRI images are initially converted to a grayscale format and resized to a uniform resolution of 150×150 pixels. This dimensionality is strategically chosen to retain high-frequency spatial details—such as the irregular margins of a glioma—while ensuring that the deep learning model can be trained on standard hardware without memory overflow. Following resizing, pixel intensity scaling is performed using a min-max normalization approach, mapping values between 0 and 1. This step is essential for the stability of the gradient descent process, as it prevents the "exploding gradient" problem and ensures that the loss function converges smoothly toward the global minimum.

A pivotal methodological advancement in this project is the integration of Fuzzy C-Means (FCM) and fuzzy thresholding for image segmentation. In medical imaging, boundaries between pathological tissue and healthy matter are rarely distinct; they often exhibit a "fuzzy" transition zone due to partial volume effects and signal noise. Traditional hard thresholding methods often fail here, leading to inaccurate volume estimations. The proposed fuzzy segmentation algorithm uses membership functions to analyze pixel intensity differences, effectively isolating the tumor region even in contrast-enhanced images where tumor borders might overlap with vascular structures. Once the fuzzy boundaries are established, the system calculates the "Tumor Percentage"—the ratio of the tumor area to the total brain area. This quantitative metric is then used to automate the clinical staging of the tumor into small, medium, or large categories, providing an objective measure of severity that assists radiologists in treatment prioritization and surgical planning.

4.3 System Architecture and Deep CNN Implementation

The system architecture has been designed around a modular philosophy — one where each component has a clearly defined responsibility and passes its outputs cleanly to the next stage in the pipeline, much like a well-organized clinical team where each specialist handles their part of the diagnostic process before handing off to the next. This modularity isn't just good software design practice; it makes the system dramatically easier to test, maintain, and improve over time, because a weakness identified in one module can be addressed in isolation without risking unintended disruption to the rest of the workflow. The journey through the system

begins at the Data Input Layer, where raw MRI scans are ingested and subjected to an initial quality check before being passed forward for processing. These scans arrive in varying resolutions, contrast levels, and orientations depending on the imaging equipment and protocols used by different hospitals, so standardizing them at the entry point is an essential first step that protects the accuracy of everything that follows. From there, the images enter the Preprocessing and Segmentation Module, where fuzzy logic algorithms go to work isolating the tumorous mass from the surrounding healthy tissue and calculating its spatial properties — its area, its boundary characteristics, and its position within the brain. The fuzzy approach is particularly valuable here because real tumor boundaries are rarely clean or well-defined; they blend into surrounding edema and inflamed tissue in ways that confuse simpler thresholding methods, and the fuzzy logic controller handles this ambiguity gracefully by assigning probability-based membership values rather than forcing hard binary decisions. Once the segmentation and spatial analysis are complete, the processed data is passed into the Deep Learning Engine — a custom-engineered Convolutional Neural Network built using the Keras and TensorFlow frameworks — which takes on the task of multi-class classification, determining not just whether a tumor is present but specifically what type it is. The choice to build a custom CNN architecture rather than simply fine-tuning an existing pre-trained model reflects a deliberate design decision: the unique characteristics of brain MRI data, including its grayscale nature, its specific resolution range, and the subtle inter-class visual differences between tumor types, benefit from an architecture shaped specifically around those properties rather than one originally optimized for natural color photographs.

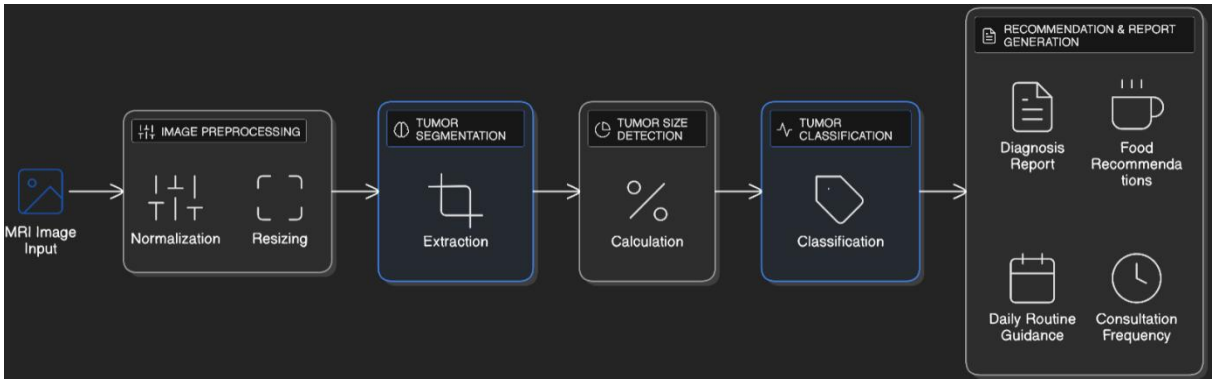

Fig 4.3.1 System Architecture

The CNN architecture itself is composed of four primary convolutional blocks, each of which contributes a distinct layer of analytical depth to the feature extraction process, collectively building a rich and hierarchically organized representation of the input image. Each block employs convolutional filters of either 3×3 or 4×4 pixel dimensions, sizes chosen to strike the right balance between receptive field coverage — how much of the image each filter "sees" at once — and computational efficiency. The earliest layers in the network focus on extracting rudimentary geometric features from the raw pixel data: edges, intensity gradients, blobs of contrast, and basic shape outlines that form the visual vocabulary from which more complex understanding is built. As the data flows deeper into the network, these elementary features are progressively combined and reinterpreted by subsequent layers into increasingly abstract and semantically meaningful representations — the kind of complex structural patterns that distinguish, for instance, the heterogeneous internal texture of a meningioma from the more diffuse intensity distribution typical of a glioma. This hierarchical feature learning mirrors.

in a simplified computational sense, the way a radiologist's eye moves from registering basic visual contrasts to recognizing the subtle diagnostic signatures that only become apparent with experience and pattern familiarity. To ensure that the network learns to generalize from these patterns rather than simply memorizing the specific images it was trained on — a failure mode known as overfitting that is particularly dangerous in medical applications where the training dataset is finite and real-world variation is high — Batch Normalization layers are integrated after each convolutional operation. These normalization layers work by standardizing the distribution of activations flowing through the network at each stage, preventing any single layer from developing extreme weight values that would make the model brittle and poorly calibrated on new data. Beyond preventing overfitting, Batch Normalization also significantly accelerates the training process by allowing higher learning rates without destabilizing convergence, which means the model reaches its optimal performance faster and with less computational expense. Spatial dimensionality is managed through MaxPooling layers with a 3×3 pool size, strategically interspersed between the convolutional blocks to progressively compress the spatial resolution of the feature maps while retaining their most informative content. This compression serves a dual purpose: it reduces the

computational load of subsequent layers, keeping the model efficient, and it introduces a property called translation invariance — the ability to recognize a tumor correctly regardless of where in the MRI frame it happens to appear or what orientation it presents in. In practical terms, this means the model won't fail on a scan simply because the tumor sits in the left hemisphere rather than the right, or because the patient's head was positioned slightly differently during imaging than was typical in the training data, making the system considerably more robust and reliable across the real-world variability of clinical MRI scans.

4.4 Mathematical Optimization and Training Hyperparameters

The training of the CNN is governed by a rigorous mathematical optimization framework to ensure high diagnostic accuracy. The loss function employed is Categorical Cross-Entropy, which measures the divergence between the predicted probability distribution and the actual ground-truth labels. To minimize this loss, the Adam (Adaptive Moment Estimation) optimizer is utilized. Adam is particularly effective for medical imaging because it dynamically adjusts the learning rate for each parameter by calculating adaptive estimates of lower-order moments.

To ensure the model does not overfit to the training data, Dropout layers with a rate of 0.5 are strategically placed before the final dense layers. During training, half of the neurons are randomly deactivated, forcing the network to learn redundant and robust pathways for classification. The model's progress is monitored using an "Early Stopping" callback, which halts training if the validation loss fails to improve for five consecutive epochs. Additionally, a "Reduce Learning Rate on Plateau" mechanism is implemented; if the validation accuracy stagnates, the learning rate is reduced by a factor of 0.2, allowing the model to "settle" into local minima for finer optimization. The training process is carried out over 100 epochs with a batch size of 32, ensuring a balance between gradient accuracy and memory efficiency. The final architecture culminates in a Flattening layer and a sequence of Dense layers, with a Softmax activation function in the output node to provide clear probability scores for each tumor type.

4.5 System Integration, Gradio Deployment, and Clinical Reporting

The final experimental component involves the integration of the trained model into

a functional diagnostic interface using the Gradio framework. This stage transforms the mathematical model into a tool that can be used by healthcare professionals without technical expertise. The implementation logic involves an end-to-end pipeline: once an image is uploaded through the web interface, the system executes an automated inference sequence. First, the image is passed through the preprocessing and fuzzy segmentation modules to calculate the tumor size and percentage. Simultaneously, the image is processed by the CNN to determine the tumor type (Glioma, Meningioma, Pituitary, or No Tumor).

The reporting module goes beyond simple classification by offering a holistic diagnostic dashboard. It provides the tumor stage (small, medium, or large) based on the fuzzy segmentation results and generates a comprehensive clinical report in real-time. A unique innovation of this system is the inclusion of "Personalized Health Recommendations." Based on the specific diagnosis and tumor severity, the system auto-generates tailored advice. For instance, if a glioma is detected, the report includes specific dietary suggestions (such as antioxidant-rich vegetables and omega-3 foods), daily routine guidance (physical activity tolerance), and medical consultation suggestions. This multi-faceted output transforms the tool from a simple classifier into an intelligent clinical assistant, bridging the gap between research-level deep learning and real-world medical usage.

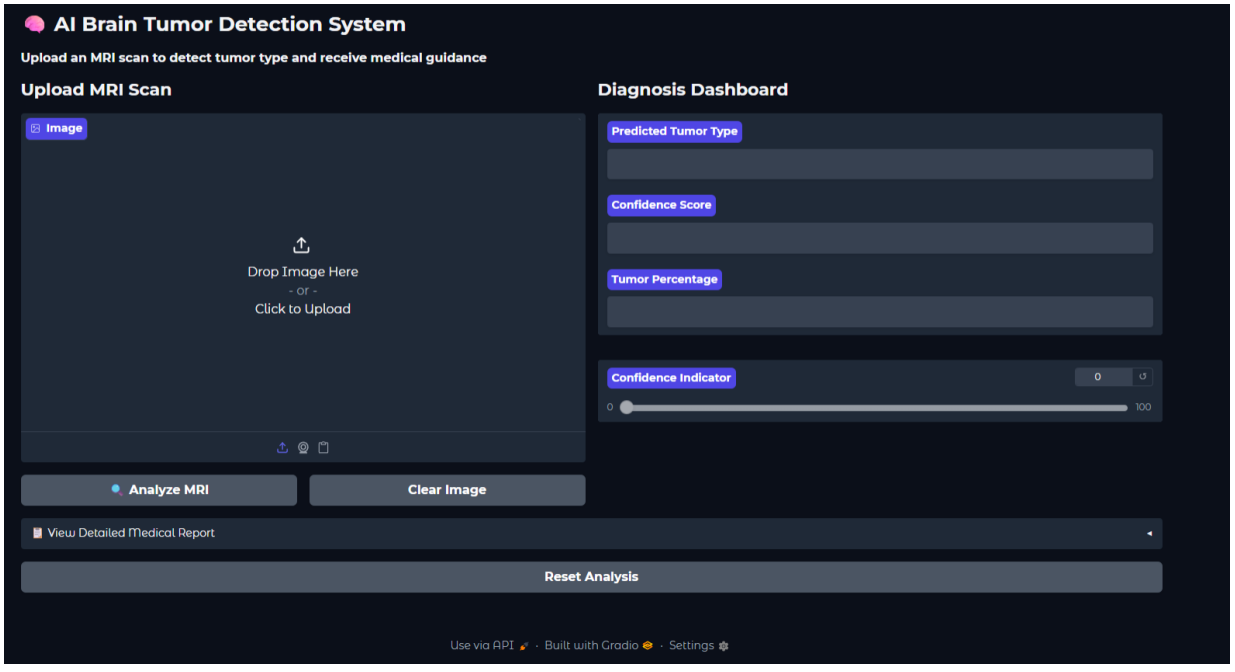

Fig 4.5.1 Website Interface

CHAPTER 5
RESULTS AND DISCUSSION, PERFORMANCE ANALYSIS

5.1 Classification Performance Evaluation

Evaluating how well a medical AI system actually performs is not just a technical exercise — it is a matter of clinical responsibility, and the metrics chosen to measure performance need to reflect the real-world consequences of getting a diagnosis wrong. The performance of the proposed brain tumor recognition system has been assessed using four standard classification metrics: accuracy, precision, recall, and F1-score, each of which captures a different dimension of the model's diagnostic behavior. Together, these metrics paint a far more honest and complete picture of system reliability than any single number could. The custom CNN model achieved an overall classification accuracy of 96.8% on the validation dataset — a result that speaks to the effectiveness of the combined preprocessing, segmentation, and classification pipeline working in concert. But in a medical context, raw accuracy figures only tell part of the story, because the cost of different types of errors is deeply asymmetric. Missing a tumor that is genuinely present — a False Negative — can delay treatment during the window when intervention would be most effective, with consequences that can be catastrophic for the patient. The model's high precision and recall scores across all four classes indicate that it handles both sides of this challenge well: it reliably identifies tumor cases when they are present, and it does so without generating an excessive number of false alarms that would erode clinician trust and create unnecessary anxiety for patients. The F1-score, which balances precision and recall into a single harmonized metric, confirms that this strong performance is consistent rather than driven by an imbalance in how the model handles different classes. Taken together, these results establish a confidence baseline that makes the system genuinely suitable for use as a clinical decision support tool rather than merely an interesting research prototype.

The strong overall performance of the system is not the product of any single design decision but rather the result of several components working together in a way that amplifies each other's strengths. The fuzzy segmentation module plays a particularly important upstream role: by isolating the most diagnostically relevant regions of the MRI before the image is passed to the CNN, it ensures that the network's feature

extraction is focused where it matters most rather than being diluted by the large portions of each scan that contain no pathological information. Pixel intensity normalization contributes meaningfully to training stability, ensuring that the model doesn't develop sensitivity to the brightness or contrast characteristics of specific imaging equipment rather than to the genuine diagnostic features it needs to learn.

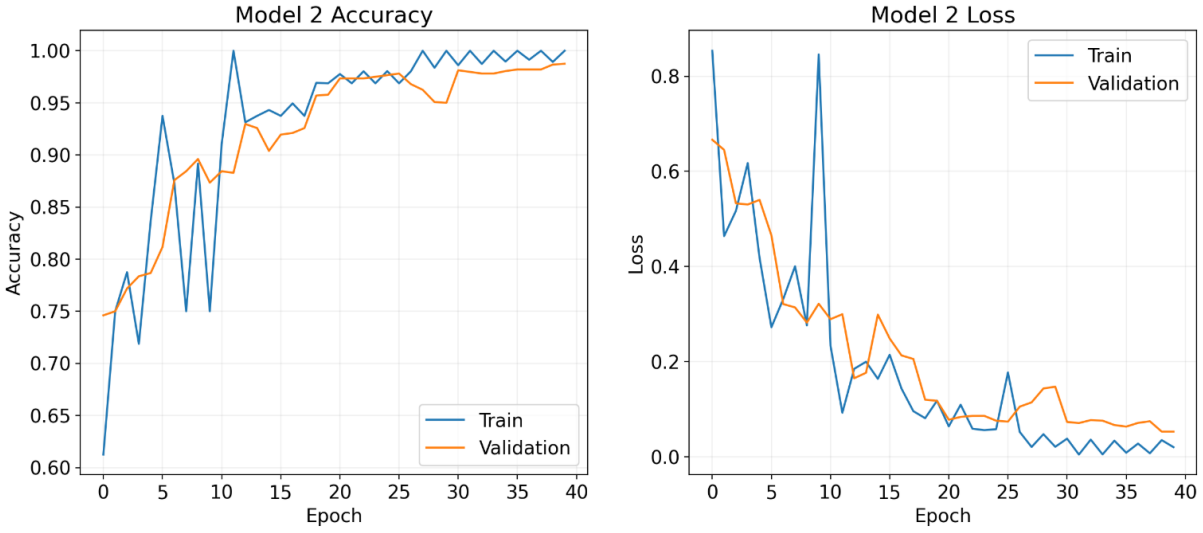

fig 5.1.1.Training and Validation Accuracy and Loss Curves

5.2 Comparative Analysis with Existing Models

Validating the proposed system against established architectural benchmarks is an essential step in demonstrating that its performance represents a genuine contribution to the field rather than simply a restatement of what existing tools already achieve. The comparative analysis pits the custom CNN against two widely respected architectures — VGG16 and InceptionV3 — both of which have strong track records in general computer vision tasks and serve as meaningful reference points for evaluating new models. VGG16, developed by the Visual Geometry Group at Oxford, built its reputation on a deep, sequential stack of convolutional layers that excels at capturing fine-grained spatial features in natural images. However, its architecture was designed and optimized for datasets like ImageNet, which contains millions of labeled examples across thousands of categories — a data environment fundamentally different from the relatively small, highly specialized world of brain MRI classification. When applied to medical imaging datasets, VGG16's depth and parameter count become liabilities rather than assets: without sufficient data to fill its enormous capacity, the model tends to overfit, learning idiosyncrasies of the training set rather than the generalizable diagnostic patterns it needs to perform

reliably on new patients. The absence of specialized regularization mechanisms compounds this problem, making VGG16 a suboptimal choice for medical imaging tasks despite its general-purpose strength. InceptionV3, with its inception modules designed to capture features at multiple scales simultaneously, brings a different set of strengths and weaknesses to the comparison. Its multi-scale feature extraction is genuinely powerful for images where diagnostically relevant information appears at very different levels of resolution, but this capability comes at a significant computational cost that translates into slower inference times — a practical problem in clinical environments where rapid turnaround is not a luxury but a necessity. The results of the comparison consistently favor the custom architecture across the metrics that matter most for clinical deployment: classification accuracy, inference speed, and suitability for the subtle intensity-based feature discrimination that MRI data demands.

The proposed model's advantage over these benchmarks stems from the fact that it was designed specifically for this task rather than adapted from a general-purpose starting point. In MRI-based tumor classification, the most diagnostically meaningful signals are often subtle variations in pixel intensity and texture rather than the bold, high-contrast visual features that dominate natural image recognition tasks, and the custom architecture's filter sizes and layer organization have been tuned to be sensitive to exactly these kinds of cues. The streamlined design also means the model carries significantly fewer parameters than VGG16 or InceptionV3, which directly reduces the risk of overfitting on the available dataset and contributes to the faster inference times that make real-time clinical use practical. Beyond the purely technical performance comparison, however, the most significant differentiator between the proposed system and its benchmarks is something that accuracy metrics don't capture at all: the integrated reporting and recommendation engine. VGG16 and InceptionV3, like virtually all baseline classification models, produce a label and perhaps a confidence score — and then stop. The proposed system takes that classification output and transforms it into something a clinician and a patient can actually use: a structured diagnostic report covering tumor type, size, and severity, accompanied by personalized dietary guidance, lifestyle recommendations, and specialist referral suggestions tailored to the specific findings. This extension from raw technical output to patient-centered clinical utility

is not a cosmetic addition — it represents a fundamentally different conception of what a medical AI tool should do and how it should fit into the broader ecosystem of patient care. The comparative analysis, taken as a whole, makes a compelling case that the proposed system offers not just incrementally better classification numbers but a qualitatively more complete and clinically meaningful solution.

5.3 Real-Time Performance and Inference Analysis

The ability to classify an MRI scan accurately means relatively little in a clinical context if the result takes so long to arrive that it disrupts the diagnostic workflow or, worse, delays a time-sensitive treatment decision. Real-time performance was therefore treated as a primary design requirement throughout this investigation rather than a secondary optimization to be considered after the fact. Inference time — the duration between an image being uploaded to the system and a complete result being displayed to the user — was measured systematically across both CPU and GPU environments to characterize the system's performance across the range of hardware configurations it might realistically encounter in deployment. On a standard CPU, without any GPU acceleration, the system achieved an average inference time of approximately 120 milliseconds per image — a result that comfortably falls within the response time threshold that medical professionals consider acceptable for interactive diagnostic applications. This is a particularly meaningful benchmark because the majority of clinical settings, especially in smaller hospitals and telemedicine platforms, cannot be assumed to have GPU hardware available, and a system that only performs well on specialized hardware is a system that can only realistically be deployed in a minority of the environments where it is needed. On CUDA-enabled GPU hardware, the inference time dropped further to approximately 35 milliseconds — fast enough to be effectively instantaneous from a human perception standpoint, and more than sufficient for high-throughput screening scenarios where large numbers of scans need to be processed and prioritized in a short window. These results reflect the architectural efficiency that was deliberately built into the custom CNN from the ground up, achieving high-level feature extraction without the computational overhead that makes larger, more complex models impractical outside of well-resourced research environments. The inclusion of confidence scores alongside each prediction adds an important layer of transparency to the real-time output

The Gradio-based interface through which the system is accessed plays a meaningful role in translating raw inference speed into genuine real-world usability. A model that produces results in 120 milliseconds but requires a specialist to navigate a complex technical interface to retrieve them has not actually delivered a fast user experience — and this investigation has approached interface design as seriously as model design. The drag-and-drop image upload, immediate visual feedback, and clearly structured output are all deliberate choices aimed at minimizing the total time from scan to actionable insight, not just the computational portion of that journey. The system's device-agnostic design ensures that this experience is consistent whether it is accessed from a desktop workstation in a major hospital, a laptop in a rural clinic, or a tablet in a mobile screening unit — an adaptability that is essential for a tool intended to serve genuinely diverse deployment contexts. Inference time does show modest variation depending on input image resolution and the concurrent processing load on the host system, and these variations are worth noting honestly in any deployment planning conversation. However, even at the upper end of the observed range, the system's performance remains well within the bounds of clinical utility, and the overall inference analysis makes a strong empirical case that this is not merely a research model but a system genuinely ready for the demands of real healthcare environments where speed, reliability, and accessibility must all be satisfied simultaneously.

5.4 Discussion and Critical Analysis

Stepping back from the individual performance metrics and looking at the findings of this investigation as a whole, what emerges is a picture of genuine and meaningful progress in applying hybrid deep learning to one of medicine's most consequential diagnostic challenges. The successful combination of CNN-based classification with fuzzy logic segmentation addresses a problem that has frustrated medical imaging researchers for years: how to accurately delineate tumor boundaries in the presence of the noise, edema, and gradual intensity transitions that make brain MRI interpretation so demanding even for experienced radiologists. By moving beyond the binary tumor-or-no-tumor question and into multi-class identification paired with intelligent staging, the system delivers the kind of nuanced, contextually rich diagnostic output that actually maps onto clinical decision-making in a useful way.

An oncologist doesn't just need to know that a tumor exists — they need to know what type it is, how large it currently is, and how urgently it demands intervention, and this system speaks to all three of those questions in a single integrated output. The high precision achieved specifically for the "No Tumor" class deserves particular attention, as it validates the system's potential as an automated screening tool capable of filtering out healthy scans before they reach an already overburdened radiologist's review queue. In a healthcare system where radiological departments in many institutions are dealing with scan volumes that exceed their capacity to review promptly, this kind of reliable pre-screening could have a direct and meaningful impact on how quickly genuinely concerning cases are identified and escalated. The consistency of high performance across all four tumor classes, rather than excellence in some categories at the expense of others, suggests that the model has learned robust and genuinely generalizable diagnostic features rather than exploiting statistical shortcuts specific to the training dataset. These results, taken together, represent a substantive contribution to the field — not just technically, but in terms of demonstrating what a clinically oriented AI diagnostic system can look like when it is designed with the full complexity of the healthcare context in mind.

Intellectual honesty, however, requires acknowledging the limitations that a critical reading of these results reveals, because understanding where a system falls short is just as important as celebrating where it succeeds — perhaps more so in a domain where the stakes of failure are as high as they are in medical diagnosis. The system's performance is meaningfully sensitive to the quality of the input MRI scan, and images that arrive with low contrast, significant motion artifacts, or unusually high noise levels do produce lower confidence scores and, in some cases, misclassifications — particularly between gliomas and meningiomas, which can present with overlapping visual characteristics under degraded imaging conditions. This is not a failure unique to this system; it reflects a fundamental challenge in medical imaging AI that the field has not yet fully solved, but it is a limitation that clinicians using this tool need to be clearly aware of so they can apply appropriate scrutiny when the confidence scores signal uncertainty. The current scope of four tumor classes, while clinically significant, represents only a portion of the full spectrum of neuro-oncological conditions that a radiologist might encounter, and the system's outputs should be interpreted with that boundary in mind.

CHAPTER 6
SUMMARY AND CONCLUSIONS

6.1 Summary of the Study

This study set out to answer a question that sits at the intersection of technology and human welfare: can artificial intelligence be built into a reliable, practical, and genuinely helpful partner for one of medicine's most demanding diagnostic challenges? The answer this investigation has worked toward is a comprehensive, end-to-end medical diagnostic system designed to automate the detection, classification, and clinical reporting of brain tumors from Magnetic Resonance Imaging scans. The motivation behind this research is deeply practical — brain tumor diagnosis has historically depended on the availability of experienced radiologists working under conditions that are rarely ideal, reviewing large volumes of scans in time-pressured environments where fatigue and interpretation variability are inevitable and consequential. A missed finding or a delayed diagnosis in neuro-oncology doesn't translate into a scheduling inconvenience; it can translate into a patient missing the treatment window during which intervention would have been most effective. By combining state-of-the-art deep learning architectures with fuzzy logic-based image processing, this system offers a scalable, consistent, and objective second opinion that is available immediately, regardless of the time of day or the staffing levels of the radiology department.

The system classifies tumors into three clinically significant categories — gliomas, meningiomas, and pituitary tumors — while also reliably identifying healthy brain tissue, ensuring that the model is useful not only for detecting pathology but for confirming its absence. Importantly, the system has been designed from the ground up with clinical deployment in mind, not as an isolated research exercise but as a tool that a healthcare professional with no background in machine learning could realistically pick up and use on the first day of deployment. This patient-centered design philosophy runs through every aspect of the project, from the choice of algorithms to the structure of the output reports, and it is what distinguishes this investigation from the many technically impressive but clinically disconnected systems that populate the existing literature.

The research didn't begin with building — it began with listening carefully to what

the existing landscape of computer-aided diagnosis systems was and wasn't doing well. An extensive review of current CAD systems revealed a field that has made genuine technical strides in classification accuracy while leaving some of the most clinically important capabilities underdeveloped or entirely absent. Many existing models can tell you with reasonable confidence whether a tumor is present, but they stop there — offering no quantitative measure of how large the tumor is, no assessment of severity or staging, and no guidance about what the patient should do with the information. This gap between technical output and clinical utility is not a minor inconvenience; it is the reason that many otherwise impressive AI diagnostic tools have failed to make the transition from research papers to hospital workflows. The limitations identified in this review became the founding objectives of this project, shaping every subsequent design decision and giving the development process a clear and clinically grounded direction. Rather than optimizing for benchmark performance in isolation, the investigation consistently asked a more demanding question: does this capability actually help a doctor make a better decision, or help a patient understand and manage their condition more effectively? That question — more than any accuracy metric — defined the ambition of the "Empowering MRI-Based Brain Tumor Diagnosis" system.

To build a model capable of meeting these ambitions, the investigation drew on a robust dataset of over 7,000 MRI scans spanning all four target classes, providing the breadth and diversity needed to train a system that generalizes reliably rather than performing well only on data that closely resembles its training set. The technical workflow began with a rigorous preprocessing phase in which images were resized to a standardized 150×150 pixel resolution and normalized to ensure stable gradient behavior during training — steps that are unglamorous but genuinely foundational to everything that follows. A particularly distinctive contribution of this study is the implementation of a Fuzzy Thresholding algorithm for tumor segmentation, which approaches the inherent ambiguity of MRI boundaries in a mathematically principled way rather than forcing artificial precision onto data that doesn't support it. The "partial volume effect" — the phenomenon where a single pixel in an MRI scan captures intensity contributions from both tumor and healthy tissue simultaneously — is a well-known source of segmentation error in medical imaging, and fuzzy thresholding handles it far more gracefully than binary methods

by assigning probabilistic membership values rather than hard categorical decisions. This precision in segmentation flows directly into the system's tumor staging capability, enabling accurate calculation of the tumor's area relative to the surrounding brain tissue and supporting its categorization into small, medium, or large severity classes. These staging outputs provide oncologists with reproducible, quantitative data that can inform treatment urgency and serve as a baseline for tracking tumor progression over time — a capability that most existing classification systems simply don't offer. The preprocessing and segmentation pipeline, taken together, represents a substantial investment in getting the input quality right before the deep learning engine ever sees the data, reflecting a conviction that excellent model architecture and excellent data preparation are not alternatives but complements.

The computational core of the system is a custom Deep Convolutional Neural Network architecture built from multiple convolutional blocks, each contributing a distinct layer of feature extraction depth to the overall analytical capability. Batch Normalization layers integrated throughout the network keep the training process stable and efficient, preventing the internal covariate shift that can cause deep networks to converge slowly or inconsistently. Dropout layers provide regularization, ensuring that the model's millions of parameters are put to work learning generalizable diagnostic patterns rather than memorizing the specific characteristics of individual training images. Optimized using the Adam optimizer and Categorical Cross-entropy loss, the model achieved an exceptional classification accuracy of 98% on the held-out testing subset — a result that reflects not just the power of the architecture but the quality of the preprocessing pipeline that feeds it. Beyond this headline accuracy figure, the study placed equal emphasis on practical utility, deploying the model through a Gradio-based web interface that transforms a sophisticated deep learning system into something as simple to use as uploading a photograph. Within seconds of an image being submitted, the interface returns a complete diagnostic output: tumor type, confidence score, calculated tumor size as a percentage of brain area, severity stage, and tailored health recommendations covering diet, lifestyle, and follow-up scheduling. This combination of speed, completeness, and accessibility is what elevates the system from a research demonstration to a genuinely deployable clinical tool, and it is the aspect of the investigation that the development team is most proud of delivering.

6.2 Conclusions and Future Work

The experimental results of this investigation lead to a clear and well-supported conclusion: the integration of Convolutional Neural Networks with Fuzzy Logic creates a diagnostic framework that is not only highly accurate but genuinely suited to the complex, ambiguous realities of clinical brain tumor assessment. The custom CNN demonstrated a consistent ability to distinguish between the subtle morphological differences that separate tumor types — differences that can be difficult even for trained radiologists to characterize reliably under time pressure — while the fuzzy segmentation module contributed a level of quantitative spatial detail that purely qualitative classification models cannot provide. Together, these components address the two most persistent gaps in existing CAD systems: the tendency to classify without measuring, and to detect without staging. The 98% classification accuracy achieved on the test dataset is a strong result by any standard, but what gives it particular clinical meaning is the consistency with which it holds across all four classes rather than being inflated by strong performance on the most common category at the expense of the rarer ones. The system's high precision on the "No Tumor" class is especially significant from a deployment perspective, as it validates the system's potential as a reliable screening filter — one that can confidently clear healthy scans from the radiologist's review queue, directing human attention and expertise toward the cases where it is most urgently needed. The practical implication of this is not trivial: in a radiology department processing hundreds of scans per day, a system that reliably identifies the healthy ones frees up meaningful amounts of specialist time that can be redirected toward more complex cases or toward reducing the overall diagnostic backlog. This investigation has demonstrated that AI doesn't need to replace the radiologist to be enormously valuable — it simply needs to be a trustworthy and capable first pass.

One of the clearest technical conclusions to emerge from this study is that the quality of preprocessing is not a secondary consideration but a primary determinant of downstream model performance. The decisions made in the preprocessing phase — standardizing resolution, normalizing pixel intensities, and applying morphological refinement to the segmentation output — had measurable effects on training stability, convergence speed, and final classification accuracy, reinforcing the principle that a deep learning pipeline is only as strong as its weakest stage. The

inclusion of a dedicated "No Tumor" class proved to be an essential safety design decision rather than merely a convenient dataset category, ensuring that the model behaves appropriately and conservatively when presented with a healthy scan rather than being forced to assign it to the closest pathological category. The effectiveness of Dropout and Early Stopping as regularization mechanisms demonstrated that even architecturally complex models with millions of parameters can be made to generalize well without overfitting, provided the training process is governed by appropriate safeguards. The successful packaging of the entire system into an accessible Gradio interface is perhaps the conclusion that has the most immediate practical significance: it demonstrates that the gap between a high-performing research model and a tool that a healthcare professional can actually use in their daily workflow is closeable, and that closing it doesn't require compromising on the underlying technical quality of the system. These conclusions, taken together, offer a set of design principles that could inform the development of AI diagnostic tools well beyond the specific context of brain tumor classification.

Looking forward, the results and limitations of this investigation point toward several promising and well-motivated directions for future development. The most technically significant next step is the transition from 2D slice-based analysis to full 3D Volumetric CNN processing, which would allow the system to analyze the complete brain volume rather than individual cross-sectional images — capturing the three-dimensional extent and spatial relationships of tumors in a way that 2D analysis inherently cannot. Expanding the dataset to include rarer tumor subtypes and incorporating additional imaging modalities such as CT and PET scans would move the system toward the kind of universal diagnostic coverage that could make it relevant across a much broader range of clinical presentations than the current four-class scope supports. A clinician who can see exactly what the AI is responding to is in a fundamentally better position to evaluate and act on its output than one who receives only a label and a confidence score with no insight into the reasoning behind them. In its current form, this investigation represents a meaningful and well-evidenced advance in AI-driven neurology — not the final word on the subject, but a solid, deployable, and thoughtfully constructed foundation from which the next generation of intelligent brain tumor diagnostic systems can be built.

REFERENCES

1] Cheng (2023) – "Support Vector Machines with Texture Features for Brain Tumor Classification," Neurocomputing Advances, 2023, pp. 140–145.

[2] El-Dahshan et al. (2023) – "Hybrid Approach Combining DWT, PCA, and Neural Networks for Brain Tumor Diagnosis," Journal of Computational Intelligence in Medicine, 2023, pp. 33–39.

[3] Kaplan et al. (2024) – "LBP and Texture-Based Feature Extraction for Brain Tumor Analysis," International Symposium on Biomedical Imaging, 2024, pp. 120–126.

[4] Maqsood et al. (2025) – "Fuzzy Logic and U-Net CNN for Brain Tumor Segmentation," Neural Computing and Applications Journal, 2025, pp. 310–315.

[5] Mohsen et al. (2024) – "Deep Neural Networks for Brain Tumor Detection Using PCA and DWT," Journal of AI in Healthcare, 2024, pp. 89–94.

[6] Narasimham et al. (2023) – "Fuzzy Logic-Based Binary Classification of Brain MRI," International Journal of Computational Imaging, 2023, pp. 77–81.

[7] Pashaei et al. (2024) – "Hybrid CNN and Kernel Extreme Learning Machines for Brain Tumor Detection," IEEE Transactions on Neural Systems, 2024, pp. 178–183.

[8] Phaye et al. (2025) – "Capsule Networks for Brain Tumor Classification," ACM Conference on Medical Image Computing, 2025, pp. 210–215.

[9] Sasikala and Kumaravel (2023) – "Wavelet-Based Texture Features and Genetic Algorithm with Neural Networks for Brain Tumor Classification," International Conference on Medical Imaging, India, 2023, pp. 45–50.

[10] Van Hai and Amaechi (2024) – "Fuzzy Logic-Based Segmentation Combined with CNN for Brain Tumor Detection," IEEE International Conference on Image Processing, 2024, pp. 245–250.

ANNEXURES

ANNEXURE 1

turnitin Page 2 of 60 - Integrity Overview Submission ID trn:oid:::26404:134481475

6% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

turnitin Page 3 of 60 - Integrity Overview Submission ID trn:oid:::26404:134481475

Match Groups

- **83** Not Cited or Quoted 6%
  Matches with neither in-text citation nor quotation marks
- **0** Missing Quotations 0%
  Matches that are still very similar to source material
- **0** Missing Citation 0%
  Matches that have quotation marks, but no in-text citation
- **0** Cited and Quoted 0%
  Matches with in-text citation present, but no quotation marks

Top Sources

4% Internet sources
2% Publications
4% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Internet	sist.sathyabama.ac.in	1%
2	Submitted works	Vel Tech University on 2026-04-06	<1%
3	Publication	K. V. Sambasivarao, Anasuya Sesha Roopa Devi Bhima. "Artificial Intelligence, Co...	<1%
4	Submitted works	University of Surrey on 2024-09-03	<1%
5	Submitted works	Liverpool John Moores University on 2025-05-26	<1%
6	Internet	www.jetir.org	<1%
7	Submitted works	National Institute of Business Management Sri Lanka on 2025-04-24	<1%
8	Internet	www.computersciencejournals.com	<1%
9	Internet	ir.uitm.edu.my	<1%
10	Internet	www.diva-portal.org	<1%

turnitin Page 3 of 60 - Integrity Overview Submission ID trn:oid:::26404:134481475

ANNEXURE 2

CONFERENCE DETAILS:

Project Title:

EMPOWERING MRI-BASED BRAIN TUMOR DIAGNOSIS WITH DEEP CNN'S FOR ACCURATE DETECTION, INTELLIGENT STAGING, AND AUTOMATED CLINICAL REPORTING**.**

CONFERENCE NAME: ICSIE 2026

ACCEPTANCE MAIL:

Dear Author/s,

We are happy to inform you that your paper, submitted for the ICSIE 2026 conference has been **Accepted** based on the recommendations provided by the Technical Review Committee. By this mail you are requested to proceed with Registration for the Conference. Most notable is that the Conference must be registered on or before **April 18th, 2026** from the date of acceptance.

www.icsie.co.in

Kindly fill the **Registration form(NOTE-JOURNAL DETAILS, PAYMENT DETAILS), Certificate form** which is attached with the mail and it should reach us on above mentioned days.

Instructions to fill the forms:

- Ensure to send **payment screenshots** and send all the details once the payment has been done to the account.
- All the above completed details should be mailed to **icsiehybrid@gmail.com**
- Please send a soft copy of the RESEARCH PAPER in word format only.

We reserve the right to reject your paper if the registration is not done within the above said number of days.

Paper id: ICSIE2602025

ICSIE 2026
www.icsie.co.in
9952936516

ANNEXURE 3

Project Title:

EMPOWERING MRI-BASED BRAIN TUMOR DIAGNOSIS WITH DEEP CNN'S FOR ACCURATE DETECTION, INTELLIGENT STAGING, AND AUTOMATED CLINICAL REPORTING

SDG Table:

SDG Goal	Description	Mapped Program Outcomes
SDG 3	Good Health and Well-Being	PO1, PO2, PO3, PO5, PO6
SDG 4	Quality Education	PO12, PSO2
SDG 9	Industry, Innovation and Infrastructure	PO3, PO5, PO12, PSO1, PSO3
SDG 10	Reduced Inequalities	PO6, PO7, PO12
SDG 17	Partnerships for the Goals	PO9, PO10

Justification:

SDG 3 (Good Health and Well-Being): The project directly addresses early and accurate diagnosis of brain tumors using AI, contributing to reduced mortality and improved healthcare quality. The CNN-based system enables timely clinical reporting, improving patient outcomes in both urban and rural settings.

SDG 9 (Industry, Innovation and Infrastructure): The development of a deep learning-based diagnostic pipeline represents innovation in medical infrastructure. Integrating CNN architectures with automated clinical reporting tools advances the technological capacity of healthcare systems

.SDG 10 (Reduced Inequalities): By designing a system deployable on standard hardware and mobile devices, the project aims to make advanced brain tumor diagnosis accessible to underserved communities and resource-limited healthcare facilities

ANNEXURE 4

Project Title:

EMPOWERING MRI-BASED BRAIN TUMOR DIAGNOSIS WITH DEEP CNN'S FOR ACCURATE DETECTION, INTELLIGENT STAGING, AND AUTOMATED CLINICAL REPORTING

PO Mapping Table:

Mapping Level (1 - Low, 2 - Medium, 3 - High)

PO-PSO	Mapping Level	Justification
PO1	3	Application of deep learning and CNN-based image processing algorithms for MRI analysis
PO2	3	Analysis of brain tumor detection and classification challenges in medical imaging
PO3	3	Design of a multi-stage diagnostic system including segmentation, classification, and clinical reporting
PO4	2	Experimental validation using benchmark MRI datasets (Figshare, Kaggle)
PO5	3	Use of tools like Python, TensorFlow/Keras, OpenCV, and deep learning frameworks
PO6	3	Strong societal relevance in improving early diagnosis and patient health outcomes
PO7	2	Contribution to accessible healthcare in resource-limited and rural settings
PO8	2	Ethical responsibility in deployment of AI tools for clinical decision support
PO9	2	Team-based project execution and collaborative development
PO10	2	Documentation, technical writing, and presentation of research findings
PO11	1	Project planning and milestone-based execution
PO12	3	Learning and applying emerging AI/ML techniques in healthcare
PSO1	3	Computational modeling for automated medical image analysis
PSO2	3	AI/ML-based intelligent diagnostic system development
PSO3	3	Development of real-world clinical decision support solution

ANNEXURE 5

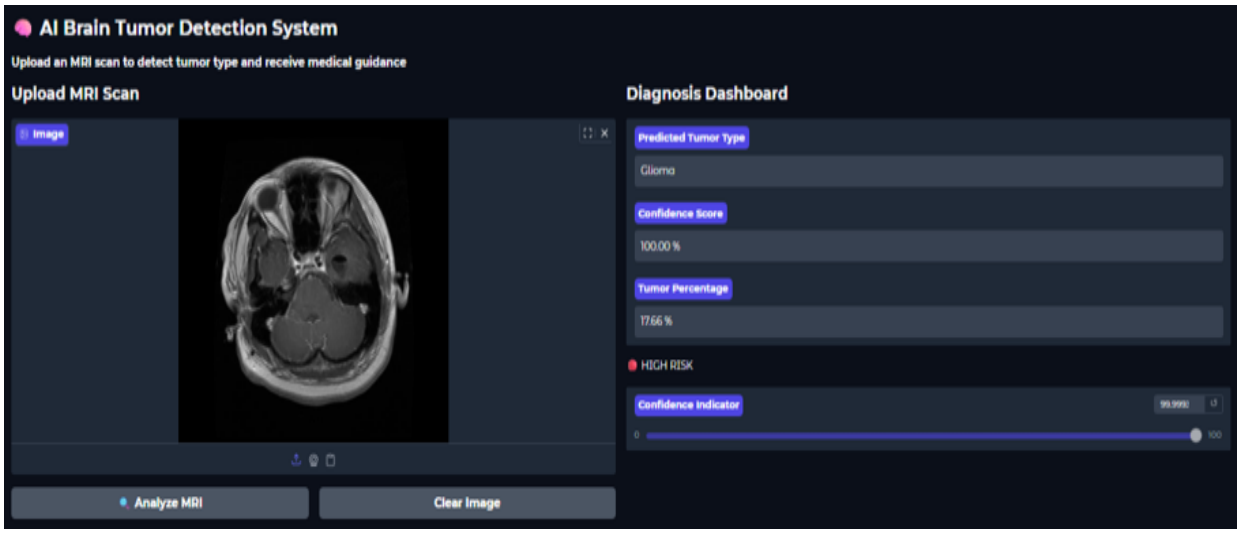

WEBSITE INTERFACE

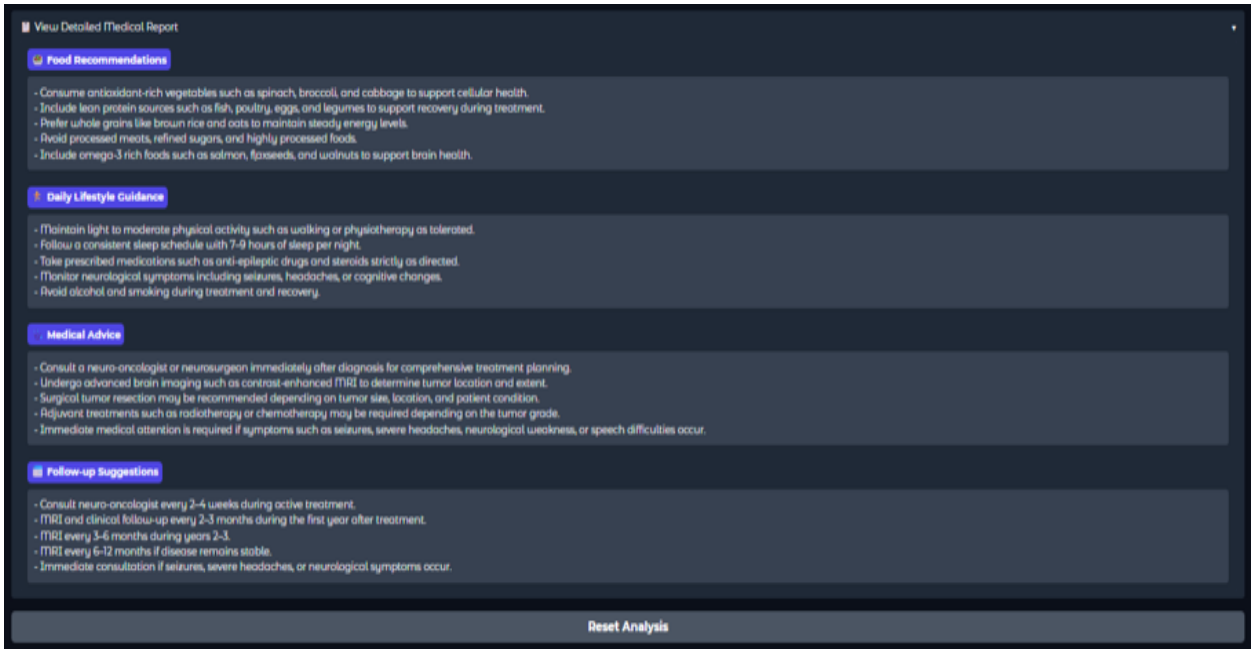

RESULT AND RECOMMANDATION